

```

#task1
import pandas as pd

# Load dataset
df = pd.read_excel("FEV-data-Excel.xlsx")

# Filter cars by budget and range
filtered_cars = df[(df['Minimal price (gross) [PLN]'] <= 350000) &
                    (df['Range (WLTP) [km]'] >= 400)]

# Group by manufacturer and calculate average battery capacity
avg_battery = filtered_cars.groupby('Make')['Battery capacity [kWh]'].mean()

# Display results
avg_battery

Make
Audi           95.000000
BMW            80.000000
Hyundai        64.000000
Kia             64.000000
Mercedes-Benz  80.000000
Tesla           68.000000
Volkswagen     70.666667
Name: Battery capacity [kWh], dtype: float64

## insights -
In task 1 Audi has the highest average battery capacity (95 kWh) among the cars that fit the customer's budget ( $\leq$  350,000 PLN) and range ( $\geq$  400 km).  

BMW and Mercedes-Benz both average 80 kWh, which is solid but lower than Audi.  

Volkswagen is around 70.7 kWh, slightly better than Tesla.  

Tesla in this filtered dataset has only 68 kWh, meaning the models under the given budget and range are not Tesla's strongest (they may have lower-range or budget versions).  

Hyundai and Kia both average the lowest at 64 kWh, which means smaller batteries compared to the premium brands.

#task2
import pandas as pd
import numpy as np
from scipy.stats import iqr
import matplotlib.pyplot as plt

##Verify the exact column names if needed
df = pd.read_excel("FEV-data-Excel.xlsx")
df.columns
col = "mean - Energy consumption [kWh/100 km]"

```

```

data=df.copy()
##Ensure the energy consumption column is numeric (coerce turns bad
strings into NaN)
data[col] = pd.to_numeric(data[col], errors="coerce")

##Compute Q1 (25th percentile) and Q3 (75th percentile)
q1 = data[col].quantile(0.25)
q3 = data[col].quantile(0.75)
iqr = q3 - q1
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr

##Boolean mask for outliers (True where value is outside the bounds)
mask_iqr = (data[col] < lower_bound) | (data[col] > upper_bound)


## Select useful columns to display for flagged rows
display_cols = [c for c in ["Car full name", "Make", "Model", col] if
c in data.columns]
outliers_iqr = data.loc[mask_iqr, display_cols].sort_values(by=col)

## Boxplot to visualize outliers
plt.figure(figsize=(6,4))
plt.boxplot(df[col], vert=False, patch_artist=True,
            boxprops=dict(facecolor="lightblue"))
plt.axvline(lower_bound, color="red", linestyle="--", label="Lower
bound")
plt.axvline(upper_bound, color="red", linestyle="--", label="Upper
bound")
plt.title("Boxplot of Energy Consumption (with Outlier Bounds)")
plt.xlabel(col)
plt.legend()
plt.show()

## Quick summary
print(f"IQR outliers found: {mask_iqr.sum()} rows")
outliers_iqr.head(10)

```

Boxplot of Energy Consumption (with Outlier Bounds)



```
IQR outliers found: 0 rows
```

```
Empty DataFrame
```

```
Columns: [Car full name, Make, Model, mean - Energy consumption  
[kWh/100 km]]
```

```
Index: []
```

```
## analysis -
```

→ In Task 2, I have applied the Interquartile Range (IQR) method to detect outliers **in** the column mean - Energy consumption [kWh/**100** km] **from** the FEV dataset.

I first ensured that the energy consumption column contained only numeric values by converting invalid entries to NaN. Then I calculated the first quartile (Q1) **and** third quartile (Q3), **and** used these to determine the IQR,

which represents the spread of the middle **50%** of the data. Based on this, I computed the lower **and** upper bounds **for** acceptable values.

Any car whose energy consumption fell below the lower bound **or** above the upper bound was flagged **as** an outlier. I then created a subset of the dataset containing only these

outliers along **with** useful details such **as** the car's name, make, model, **and** energy consumption.

Finally, I summarized the results by printing the total number of outliers detected **and** displayed the first few rows of the identified outlier vehicles.

```

## task 3
from scipy.stats import linregress, pearsonr, spearmanr
## Robustly detect the two columns we need
battery_col = next((c for c in df.columns if "Battery" in c and "kWh" in c), None)
range_col    = next((c for c in df.columns if "Range" in c and "km" in c), None)

data = df[[battery_col, range_col]].copy()
data[battery_col] = pd.to_numeric(data[battery_col], errors="coerce")
data[range_col]   = pd.to_numeric(data[range_col], errors="coerce")

clean = data.dropna().reset_index(drop=True)

# Just show the first 10 rows in Jupyter
clean.head(10)

##Compute correlation measures
pearson_r, pearson_p = pearsonr(clean[battery_col], clean[range_col])
spearman_rho, spearman_p = spearmanr(clean[battery_col],
clean[range_col])

## Fit a linear regression (for the trend line on the plot)
slope, intercept, r_value, p_value, std_err =
linregress(clean[battery_col], clean[range_col])
r_squared = r_value**2


## Print key statistics
print(f"Pearson r = {pearson_r:.4f} (p = {pearson_p:.4e})")
print(f"Spearman rho = {spearman_rho:.4f} (p = {spearman_p:.4e})")
print(f"Linear regression: slope = {slope:.4f}, intercept = {intercept:.2f}, R^2 = {r_squared:.4f}, p = {p_value:.4e}")

## Plot the scatter and regression line
plt.figure(figsize=(8,6))
plt.scatter(clean[battery_col], clean[range_col])
x_vals = np.linspace(clean[battery_col].min(),
clean[battery_col].max(), 200)
y_vals = intercept + slope * x_vals
plt.plot(x_vals, y_vals)
plt.xlabel(battery_col)
plt.ylabel(range_col)
plt.title("Battery capacity vs Range (WLTP)")
plt.grid(True)
plt.show()

```

Pearson r = 0.8104 (p = 1.9464e-13)
Spearman rho = 0.8519 (p = 6.1453e-16)

Linear regression: slope = 3.9839, intercept = 128.45, R² = 0.6568, p = 1.9464e-13



→ In Task 3, I explored how a car's battery capacity relates to its driving range. First I picked out the right columns from the dataset – one for battery size (in kWh) and one for range (in km). Then cleaned the data by converting the values to numbers and dropping any rows with missing entries, so that the analysis was only based on valid data points.

Once the data was ready, I calculated two types of correlations. The Pearson correlation measured how strongly battery size and range are related in a straight-line (linear) way, while the Spearman correlation checked whether the relationship still holds when just looking at the general ranking of values. Both helped confirm whether cars with bigger batteries tend to travel farther.

After that, I ran a linear regression, which gave me a slope, intercept, and R² value. The slope tells us how much the range increases for every extra unit of battery capacity, and the R² value

shows how well battery size explains the variation `in` driving `range`. To make the results easier to understand, I also created a scatter plot showing each car's battery capacity versus its `range`, `and` added the regression line to highlight the overall trend. This makes it clear how battery size influences driving distance `and` whether the relationship `is` strong `or` weak.

```
## task 4
import pandas as pd
import numpy as np

class EVRecommender:
    """
    EVRecommender(df)
    Recommender for EV dataset.
    """

    def __init__(self, df,
                 price_col='Minimal price (gross) [PLN]',
                 range_col='Range (WLTP) [km]',
                 batt_col='Battery capacity [kWh]'):
        # keep original
        self.df_raw = df.copy()
        self.price_col = price_col
        self.range_col = range_col
        self.batt_col = batt_col
        # clean data
        self.df = self._prepare_data(self.df_raw)

    def _prepare_data(self, df):
        """Convert to numeric and drop missing values in the 3 key
        columns."""
        df = df.copy()
        df[self.price_col] = pd.to_numeric(df[self.price_col],
                                         errors='coerce')
        df[self.range_col] = pd.to_numeric(df[self.range_col],
                                         errors='coerce')
        df[self.batt_col] = pd.to_numeric(df[self.batt_col],
                                         errors='coerce')
        df_clean = df.dropna(subset=[self.price_col, self.range_col,
                                     self.batt_col]).reset_index(drop=True)
        return df_clean

    def _score_matches(self, df_candidates, budget, desired_range,
                      desired_battery, weights):
        """
        Assigns a composite score to each candidate if strict
        filtering fails.
        """
        w_price, w_range, w_batt = weights
```

```

s = w_price + w_range + w_batt
w_price, w_range, w_batt = w_price/s, w_range/s, w_batt/s

max_price = max(1.0, df_candidates[self.price_col].max())
max_range = max(1.0, df_candidates[self.range_col].max())
max_batt = max(1.0, df_candidates[self.batt_col].max())

# price score
if budget is None or budget <= 0:
    price_score = np.zeros(len(df_candidates))
else:
    price_score = (budget - df_candidates[self.price_col]) /
budget
    price_score = price_score.clip(-1, 1)

# range score
if desired_range is None or desired_range <= 0:
    range_score = df_candidates[self.range_col] / max_range
else:
    range_score = (df_candidates[self.range_col] -
desired_range) / max(desired_range, 1)
    range_score = range_score.clip(-1, 2)

# battery score
if desired_battery is None or desired_battery <= 0:
    batt_score = df_candidates[self.batt_col] / max_batt
else:
    batt_score = (df_candidates[self.batt_col] -
desired_battery) / max(desired_battery, 1)
    batt_score = batt_score.clip(-1, 2)

composite = w_price*price_score + w_range*range_score +
w_batt*batt_score

scored = df_candidates.copy()
scored['_price_score'] = price_score
scored['_range_score'] = range_score
scored['_batt_score'] = batt_score
scored['_score'] = composite
return scored

def recommend(self, budget=None, desired_range=None,
desired_battery=None,
                      top_n=3, strict=True, weights=(0.5, 0.3, 0.2)):
"""
Recommend EVs given user preferences.
"""
df = self.df

# --- Strict filter ---

```

```

    if strict:
        mask = pd.Series(True, index=df.index)
        if budget is not None:
            mask &= (df[self.price_col] <= budget)
        if desired_range is not None:
            mask &= (df[self.range_col] >= desired_range)
        if desired_battery is not None:
            mask &= (df[self.batt_col] >= desired_battery)

        strict_matches = df[mask].copy().reset_index(drop=True)
        if len(strict_matches) >= top_n:
            strict_matches = strict_matches.sort_values(
                [self.price_col, self.range_col],
                ascending=[True, False]
            )
        return strict_matches.head(top_n)

    # --- Scoring fallback ---
    candidates = df.copy()
    scored = self._score_matches(candidates, budget,
desired_range, desired_battery, weights)
    scored_sorted = scored.sort_values('_score',
ascending=False).reset_index(drop=True)

    return scored_sorted.head(top_n)

# Load dataset
df = pd.read_excel("FEV-data-Excel.xlsx")

# Create recommender
rec = EVRecommender(df)

# Example 1: strict filter (budget=350000, range=400 km, battery=60 kWh)
print("Strict filter result:")
print(rec.recommend(budget=350000, desired_range=400,
desired_battery=60, top_n=3))

# Example 2: scoring fallback (if no strict matches found)
print("\nFallback scoring result:")
print(rec.recommend(budget=150000, desired_range=500,
desired_battery=80, top_n=3, strict=True))

Strict filter result:
      Car full name     Make           Model \
4          Kia e-Soul 64kWh     Kia       e-Soul 64kWh
3          Kia e-Niro 64kWh     Kia       e-Niro 64kWh
2  Hyundai Kona electric 64kWh  Hyundai  Kona electric 64kWh

```

	Minimal price (gross) [PLN]	Engine power [KM]	Maximum torque [Nm]
4	160990	204	395
3	167990	204	395
2	178400	204	395
	Type of brakes	Drive type	Battery capacity [kWh]
4	disc (front + rear)	2WD (front)	64.0
3	disc (front + rear)	2WD (front)	64.0
2	disc (front + rear)	2WD (front)	64.0
	Range (WLTP) [km]	...	Permissible gross weight [kg]
4	452	...	1682.0
3	455	...	2230.0
2	449	...	2170.0
	Maximum load capacity [kg]	Number of seats	Number of doors
4	498.0	5	5
3	493.0	5	5
2	485.0	5	5
	Tire size [in]	Maximum speed [kph]	Boot capacity (VDA) [l]
4	17	167	315.0
3	17	167	451.0
2	17	167	332.0
	Acceleration 0-100 kph [s]	Maximum DC charging power [kW]	
4	7.9	100	
3	7.8	100	
2	7.6	100	
	mean - Energy consumption [kWh/100 km]		
4		15.7	
3		15.9	
2		15.4	
[3 rows x 25 columns]			
Fallback scoring result:			
	Car full name	Make	Model
0	Skoda Citigo-e iV	Skoda	Citigo-e iV
1	Volkswagen ID.3 Pro S	Volkswagen	ID.3 Pro S
2	Citroën ë-C4	Citroën	ë-C4
	Minimal price (gross) [PLN]	Engine power [KM]	Maximum torque [Nm]
0	82050	83	212

```

1          179990           204           310
2          125000           136           260

      Type of brakes  Drive type  Battery capacity [kWh] \
0  disc (front) + drum (rear)  2WD (front)            36.8
1  disc (front) + drum (rear)  2WD (rear)             77.0
2        disc (front + rear)  2WD (front)            50.0

      Range (WLTP) [km]  ...  Tire size [in]  Maximum speed [kph] \
0            260  ...           14            130
1            549  ...           19            160
2            350  ...           16            150

      Boot capacity (VDA) [l]  Acceleration 0-100 kph [s] \
0            250.0                  12.3
1            385.0                  7.9
2            380.0                  9.5

      Maximum DC charging power [kW]  mean - Energy consumption [kWh/100
km] \
0                      40
15.45
1                      125
15.90
2                      100
NaN

      _price_score  _range_score  _batt_score  _score
0      0.453000    -0.480     -0.5400   -0.025500
1     -0.199933     0.098     -0.0375   -0.078067
2      0.166667    -0.300     -0.3750   -0.081667

```

[3 rows x 29 columns]

##analysis-

- In Task 4, I created an EV recommender system that suggests electric cars based on a user's budget, desired range, and battery size. The system first cleans the data, then tries a strict filter to find cars that exactly match the requirements.
- If no perfect matches are found, it falls back to a scoring method, which ranks cars by how closely they fit the preferences, using weighted scores for price, range, and battery.
- I tested it with two cases—one where strict filtering worked and another where the fallback scoring kicked in—showing that the recommender can handle both exact matches and near matches effectively.

```

## task 5 hypothesis testing between tesla and audi
## Interpretation -Hypothesis Setup
##  $H_0$  (Null Hypothesis): Tesla & Audi have the same average engine power.
##  $H_1$  (Alternative Hypothesis): Tesla & Audi have different average engine power.

import pandas as pd
from scipy.stats import ttest_ind

# Load dataset
df = pd.read_excel("FEV-data-Excel.xlsx")

# Filter engine power values for Tesla and Audi
tesla_power = df[df["Make"] == "Tesla"]["Engine power [KM]"].dropna()
audi_power = df[df["Make"] == "Audi"]["Engine power [KM]"].dropna()

# Perform two-sample independent t-test (Welch's t-test, safer when variances differ)
t_stat, p_value = ttest_ind(tesla_power, audi_power, equal_var=False)

print("T-statistic:", t_stat)
print("P-value:", p_value)

# Set significance level
alpha = 0.05
if p_value < alpha:
    print("Reject Null Hypothesis: Significant difference in average engine power between Tesla and Audi.")
else:
    print("Fail to Reject Null Hypothesis: No significant difference in average engine power between Tesla and Audi.")

```

```

T-statistic: 1.7939951827297178
P-value: 0.10684105068839565
Fail to Reject Null Hypothesis: No significant difference in average engine power between Tesla and Audi.

```

insights

- In task 5, Since p-value (**0.107**) > **0.05**, we fail to reject the null hypothesis.
- This means that although Tesla's cars have a higher average engine power (**533 KM** vs **392 KM**), the difference **is not** statistically significant at the **5%** level.
- Tesla seems to offer more powerful EVs on average, but the variation within models means the difference **is not** strong enough statistically. Customers choosing between Audi **and** Tesla should **not** focus only on

engine power – instead, factors like **range**, charging speed, **and** price may be more decisive.

Project Video Explanation :

[Click here to watch the video](#)

